

Greetings to everyone!!

- Ill be creating a model on the basis of data of the Loan Defaulters

Importing the Libraries

- Numpy
- Pandas
- Matplotlib
- Seaborn
- From sklearn
 - Scipy
 - Label Encoder
 - Model Selection
 - Decision Tree
 - Random Forest
- From imblearn
 - Random OverSampling

In [1]:

```
1 import pandas as pd
2 import numpy as np
3 import matplotlib.pyplot as plt
4 import seaborn as sns
5
6 import warnings
7 warnings.filterwarnings("ignore")
```

In [2]:

```

1 #calling the data from csv file and checking the head of the data
2 df = pd.read_csv("credit_train.csv")
3 df.head()

```

Out[2]:

	Loan ID	Customer ID	Loan Status	Current Loan Amount	Term	Credit Score	Annual Income	Years in current job	Hon Ownersh
0	14dd8831-6af5-400b-83ec-68e61888a048	981165ec-3274-42f5-a3b4-d104041a9ca9	Fully Paid	445412.0	Short Term	709.0	1167493.0	8 years	Hor Mortga
1	4771cc26-131a-45db-b5aa-537ea4ba5342	2de017a3-2e01-49cb-a581-08169e83be29	Fully Paid	262328.0	Short Term	NaN	NaN	10+ years	Hor Mortga
2	4eed4e6a-aa2f-4c91-8651-ce984ee8fb26	5efb2b2b-bf11-4dfd-a572-3761a2694725	Fully Paid	99999999.0	Short Term	741.0	2231892.0	8 years	Own Hor
3	77598f7b-32e7-4e3b-a6e5-06ba0d98fe8a	e777faab-98ae-45af-9a86-7ce5b33b1011	Fully Paid	347666.0	Long Term	721.0	806949.0	3 years	Own Hor
4	d4062e70-befa-4995-8643-a0de73938182	81536ad9-5ccf-4eb8-befb-47a4d608658e	Fully Paid	176220.0	Short Term	NaN	NaN	5 years	Re

Exploratory Data Analysis

In [3]:

```

1 # Information about the data
2 df.info()

```

1	customer_id	100000	non-null	object
2	Loan Status	100000	non-null	object
3	Current Loan Amount	100000	non-null	float64
4	Term	100000	non-null	object
5	Credit Score	80846	non-null	float64
6	Annual Income	80846	non-null	float64
7	Years in current job	95778	non-null	object
8	Home Ownership	100000	non-null	object
9	Purpose	100000	non-null	object
10	Monthly Debt	100000	non-null	float64
11	Years of Credit History	100000	non-null	float64
12	Months since last delinquent	46859	non-null	float64
13	Number of Open Accounts	100000	non-null	float64
14	Number of Credit Problems	100000	non-null	float64
15	Current Credit Balance	100000	non-null	float64
16	Maximum Open Credit	99998	non-null	float64
17	Bankruptcies	99796	non-null	float64
18	Tax Liens	99990	non-null	float64

dtypes: float64(12), object(7)
memory usage: 14.6+ MB

In [4]:

```

1 # Descriptive analysis of the Data
2 df.describe()

```

Out[4]:

	Current Loan Amount	Credit Score	Annual Income	Monthly Debt	Years of Credit History	Months since last delinquent
count	1.000000e+05	80846.000000	8.084600e+04	100000.000000	100000.000000	46859.000000
mean	1.176045e+07	1076.456089	1.378277e+06	18472.412336	18.199141	34.901321
std	3.178394e+07	1475.403791	1.081360e+06	12174.992609	7.015324	21.997829
min	1.080200e+04	585.000000	7.662700e+04	0.000000	3.600000	0.000000
25%	1.796520e+05	705.000000	8.488440e+05	10214.162500	13.500000	16.000000
50%	3.122460e+05	724.000000	1.174162e+06	16220.300000	16.900000	32.000000
75%	5.249420e+05	741.000000	1.650663e+06	24012.057500	21.700000	51.000000
max	1.000000e+08	7510.000000	1.655574e+08	435843.280000	70.500000	176.000000

By looking at the above description we can say that

- The mean value of monthly Debt is 18742.4
- Maximum Credit problems arises are 15

In [5]:

```
1 #As the Loan ID and Customer ID column is not needed, so I dropped the columns
2
3 df.drop(["Loan ID", "Customer ID"], axis=1, inplace = True)
```

In [6]:

```
1 df.isna().sum()
2
3 #To check the missing values in the Dataset
```

Out[6]:

Loan Status	514
Current Loan Amount	514
Term	514
Credit Score	19668
Annual Income	19668
Years in current job	4736
Home Ownership	514
Purpose	514
Monthly Debt	514
Years of Credit History	514
Months since last delinquent	53655
Number of Open Accounts	514
Number of Credit Problems	514
Current Credit Balance	514
Maximum Open Credit	516
Bankruptcies	718
Tax Liens	524


dtype: int64

In [7]:

```
1 #To visualize the null values in the Data set
2 sns.heatmap(df.isnull(), yticklabels=False, cbar=False, cmap="viridis")
```

Out[7]:

<AxesSubplot:>



In [8]:

```
1 #Looking at the dataset, columnns which have more than 50% missing values needs to be c
2
3 df.drop(["Months since last delinquent"],axis = 1, inplace = True)
4 df.drop(["Credit Score"],axis = 1, inplace = True)
5 df.drop(["Annual Income"],axis = 1, inplace = True)
```

In [9]:

```
1 # Coulmns that has less than 50% missing values were simply handled by dropna feature
2 df.dropna(inplace = True)
```

In [10]:

```
1 df.head()
```

Out[10]:

	Loan Status	Current Loan Amount	Term	Years in current job	Home Ownership	Purpose	Monthly Debt	Years of Credit History	Number of Open Accounts	F
0	Fully Paid	445412.0	Short Term	8 years	Home Mortgage	Home Improvements	5214.74	17.2	6.0	
1	Fully Paid	262328.0	Short Term	10+ years	Home Mortgage	Debt Consolidation	33295.98	21.1	35.0	
2	Fully Paid	9999999.0	Short Term	8 years	Own Home	Debt Consolidation	29200.53	14.9	18.0	
3	Fully Paid	347666.0	Long Term	3 years	Own Home	Debt Consolidation	8741.90	12.0	9.0	
4	Fully Paid	176220.0	Short Term	5 years	Rent	Debt Consolidation	20639.70	6.1	15.0	

In [11]:

```
1 df.isna().sum()
```

Out[11]:


```
Loan Status      0
Current Loan Amount  0
Term             0
Years in current job  0
Home Ownership   0
Purpose          0
Monthly Debt     0
Years of Credit History  0
Number of Open Accounts  0
Number of Credit Problems  0
Current Credit Balance  0
Maximum Open Credit  0
Bankruptcies     0
Tax Liens        0
dtype: int64
```

In [12]:

```
1 sns.heatmap(df.isnull(), yticklabels=False, cbar=False, cmap="viridis")
```

Out[12]:

<AxesSubplot:>



Now there are no missing values!!

In [13]:

1 df.info()

<class 'pandas.core.frame.DataFrame'>

Int64Index: 95572 entries, 0 to 99998

Data columns (total 14 columns):

#	Column	Non-Null Count	Dtype
0	Loan Status	95572 non-null	object
1	Current Loan Amount	95572 non-null	float64
2	Term	95572 non-null	object
3	Years in current job	95572 non-null	object
4	Home Ownership	95572 non-null	object
5	Purpose	95572 non-null	object
6	Monthly Debt	95572 non-null	float64
7	Years of Credit History	95572 non-null	float64
8	Number of Open Accounts	95572 non-null	float64
9	Number of Credit Problems	95572 non-null	float64
10	Current Credit Balance	95572 non-null	float64
11	Maximum Open Credit	95572 non-null	float64
12	Bankruptcies	95572 non-null	float64
13	Tax Liens	95572 non-null	float64

dtypes: float64(9), object(5)

memory usage: 10.9+ MB

In [14]:

```
1 # checking the correlation between the columns
2 df.corr().style.background_gradient()
```

Out[14]:

	Current Loan Amount	Monthly Debt	Years of Credit History	Number of Open Accounts	Number of Credit Problems	Current Credit Balance	Maximum Open Credit	Bankruptcies
Current Loan Amount	1.000000	-0.007967	0.018941	0.000731	-0.003692	0.003870	-0.001092	-0.
Monthly Debt	-0.007967	1.000000	0.218537	0.407513	-0.049283	0.482037	0.039326	-0.
Years of Credit History	0.018941	0.218537	1.000000	0.139674	0.057824	0.215992	0.030196	0.
Number of Open Accounts	0.000731	0.407513	0.139674	1.000000	-0.009716	0.225043	0.030088	-0.
Number of Credit Problems	-0.003692	-0.049283	0.057824	-0.009716	1.000000	-0.110298	-0.011487	0.
Current Credit Balance	0.003870	0.482037	0.215992	0.225043	-0.110298	1.000000	0.136555	-0.
Maximum Open Credit	-0.001092	0.039326	0.030196	0.030088	-0.011487	0.136555	1.000000	-0.
Bankruptcies	-0.001133	-0.072196	0.061097	-0.017804	0.751056	-0.119691	-0.013958	1.
Tax Liens	-0.003228	0.021759	0.017175	0.005862	0.581465	-0.015566	-0.001068	0.

- By this analysis we can say that,

- 1) Bankruptcies and Tax liens has high correlation with Number of credit problems
- 2) Current credit balance has some correlation with Monthly debt

In [15]:


```
1 df["Loan Status"].value_counts()
```

Out[15]:

```
Fully Paid      74257  
Charged Off     21315  
Name: Loan Status, dtype: int64
```

In [16]:

```
1 # checking the count of people who have fully paid and failed to pay  
2 sns.countplot(data = df, x= "Loan Status")  
3 plt.grid(True)
```



In [17]:


```
1 df["Purpose"].value_counts()
```

Out[17]:

```
Debt Consolidation    75277
other                 5665
Home Improvements    5511
Other                 3104
Business Loan         1519
Buy a Car             1208
Medical Bills         1041
Buy House             652
Take a Trip           546
major_purchase        333
small_business        272
moving                139
wedding               111
vacation              95
Educational Expenses  89
renewable_energy      10
Name: Purpose, dtype: int64
```

In [18]:

```
1 # Purpose of the Loan that was taken
2 sns.countplot(data = df, x= "Purpose")
3 plt.title("Purpose of Loan")
4 plt.xlabel("Purpose")
5 plt.ylabel("Applicant")
6 plt.xticks(rotation=90)
7 plt.grid(True)
```



By this visualization we can say that

- the highest reason for the loan was Debt consolidation

In [19]:

```
1 df["Years in current job"].value_counts()
```


Out[19]:

10+ years	31090
2 years	9104
3 years	8151
< 1 year	8114
5 years	6778
1 year	6436
4 years	6132
6 years	5676
7 years	5573
8 years	4569
9 years	3949

Name: Years in current job, dtype: int64

In [20]:

```
1 sns.countplot(data = df, x= "Years in current job")
2 plt.title("Years of Experience")
3 plt.xlabel("Years")
4 plt.ylabel("Applicant")
5 plt.xticks(rotation=45)
6 plt.show()
7
```

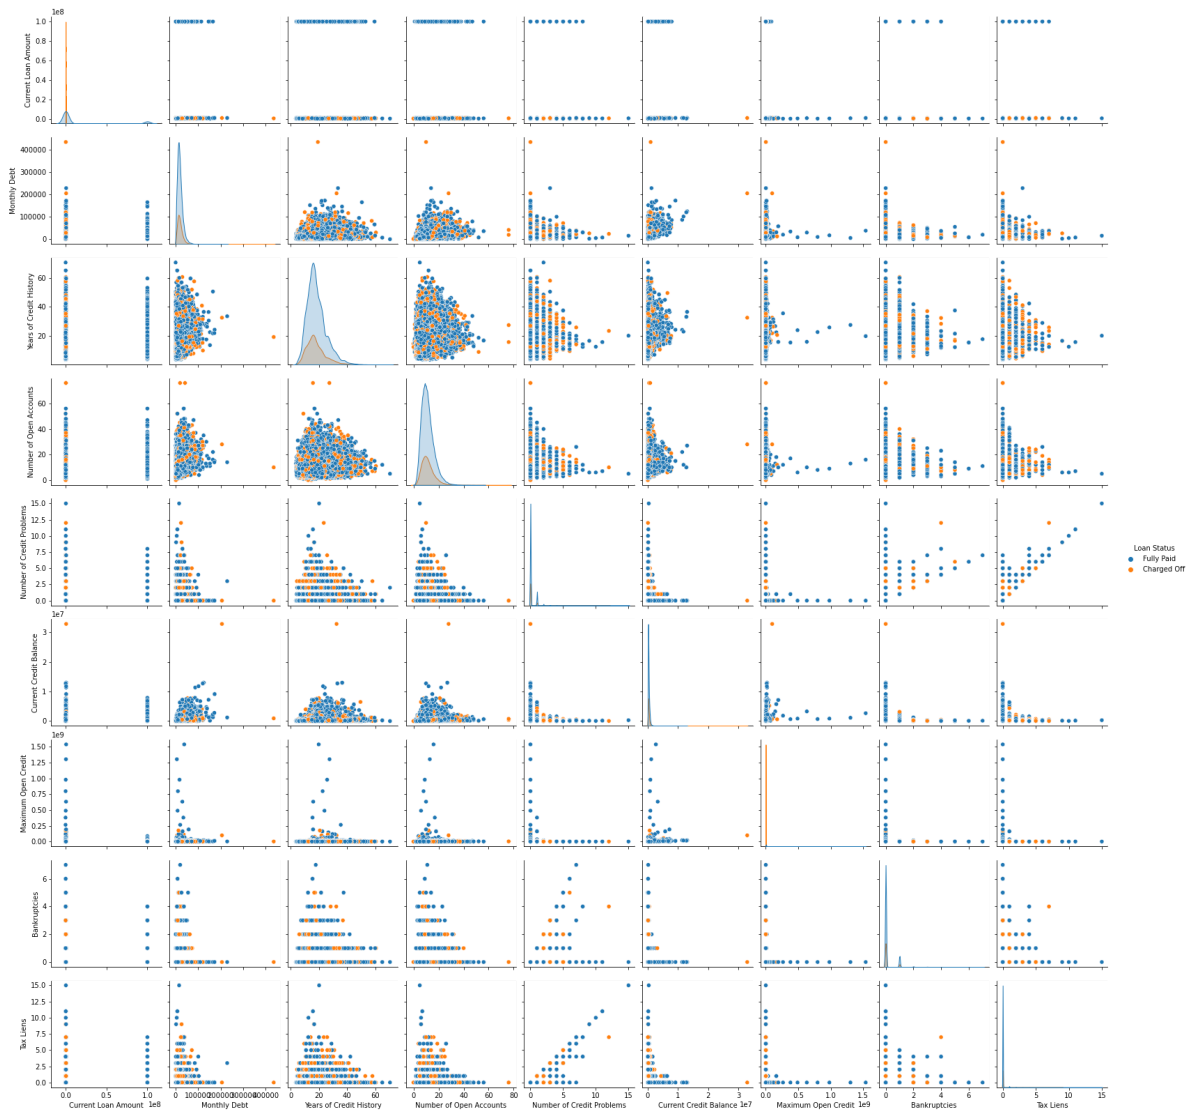


In [21]:

```
1 # checking the pairplot for the Dataset
2 sns.pairplot(df, hue="Loan Status")
```

Out[21]:

<seaborn.axisgrid.PairGrid at 0x2bf0532c10>



In [22]:

```
1 df.head()
```

Out[22]:

	Loan Status	Current Loan Amount	Term	Years in current job	Home Ownership	Purpose	Monthly Debt	Years of Credit History	Number of Open Accounts	F
0	Fully Paid	445412.0	Short Term	8 years	Home Mortgage	Home Improvements	5214.74	17.2	6.0	
1	Fully Paid	262328.0	Short Term	10+ years	Home Mortgage	Debt Consolidation	33295.98	21.1	35.0	
2	Fully Paid	99999999.0	Short Term	8 years	Own Home	Debt Consolidation	29200.53	14.9	18.0	
3	Fully Paid	347666.0	Long Term	3 years	Own Home	Debt Consolidation	8741.90	12.0	9.0	
4	Fully Paid	176220.0	Short Term	5 years	Rent	Debt Consolidation	20639.70	6.1	15.0	

Data preprocessing

In [23]:

```
1 # splitting the data into X and Y
2 x = df.iloc[:,1:]
3 y = df.iloc[:,1]
```

In [24]:

1 x

Out[24]:

	Current Loan Amount	Term	Years in current job	Home Ownership	Purpose	Monthly Debt	Years of Credit History	Number of Open Accounts	Nu of C Prob
0	445412.0	Short Term	8 years	Home Mortgage	Home Improvements	5214.74	17.2	6.0	
1	262328.0	Short Term	10+ years	Home Mortgage	Debt Consolidation	33295.98	21.1	35.0	
2	9999999.0	Short Term	8 years	Own Home	Debt Consolidation	29200.53	14.9	18.0	
3	347666.0	Long Term	3 years	Own Home	Debt Consolidation	8741.90	12.0	9.0	
4	176220.0	Short Term	5 years	Rent	Debt Consolidation	20639.70	6.1	15.0	
...	
99994	210584.0	Short Term	1 year	Home Mortgage	Other	3727.61	17.4	6.0	
99995	147070.0	Short Term	7 years	Own Home	other	2202.86	22.3	5.0	
99996	9999999.0	Short Term	1 year	Rent	Debt Consolidation	13109.05	9.4	22.0	
99997	103136.0	Short Term	6 years	Rent	Debt Consolidation	7315.57	18.8	12.0	
99998	530332.0	Short Term	9 years	Rent	Debt Consolidation	9890.07	15.0	8.0	

95572 rows × 13 columns



In [25]:

1

y

Loan Status	
0	Fully Paid
1	Fully Paid
2	Fully Paid
3	Fully Paid
4	Fully Paid
...	...
99994	Fully Paid
99995	Fully Paid
99996	Fully Paid
99997	Fully Paid
99998	Fully Paid

In [26]:

1

df_cat = df.select_dtypes(object)

2

df_num = df.select_dtypes(["int64", "float64"])

In [27]:

1

df_cat

Out[27]:

Loan Status		Term	Years in current job	Home Ownership	Purpose
0	Fully Paid	Short Term	8 years	Home Mortgage	Home Improvements
1	Fully Paid	Short Term	10+ years	Home Mortgage	Debt Consolidation
2	Fully Paid	Short Term	8 years	Own Home	Debt Consolidation
3	Fully Paid	Long Term	3 years	Own Home	Debt Consolidation
4	Fully Paid	Short Term	5 years	Rent	Debt Consolidation
...
99994	Fully Paid	Short Term	1 year	Home Mortgage	Other
99995	Fully Paid	Short Term	7 years	Own Home	other
99996	Fully Paid	Short Term	1 year	Rent	Debt Consolidation
99997	Fully Paid	Short Term	6 years	Rent	Debt Consolidation
99998	Fully Paid	Short Term	9 years	Rent	Debt Consolidation

95572 rows × 5 columns

In [28]:

1 df_num

Out[28]:

	Current Loan Amount	Monthly Debt	Years of Credit History	Number of Open Accounts	Number of Credit Problems	Current Credit Balance	Maximum Open Credit	Bankruptcies
0	445412.0	5214.74	17.2	6.0	1.0	228190.0	416746.0	1.0
1	262328.0	33295.98	21.1	35.0	0.0	229976.0	850784.0	0.0
2	99999999.0	29200.53	14.9	18.0	1.0	297996.0	750090.0	0.0
3	347666.0	8741.90	12.0	9.0	0.0	256329.0	386958.0	0.0
4	176220.0	20639.70	6.1	15.0	0.0	253460.0	427174.0	0.0
...
99994	210584.0	3727.61	17.4	6.0	0.0	456.0	259160.0	0.0
99995	147070.0	2202.86	22.3	5.0	0.0	47766.0	658548.0	0.0
99996	99999999.0	13109.05	9.4	22.0	0.0	153045.0	509234.0	0.0
99997	103136.0	7315.57	18.8	12.0	1.0	109554.0	537548.0	1.0
99998	530332.0	9890.07	15.0	8.0	0.0	404225.0	738254.0	0.0

95572 rows × 9 columns

In [29]:

1 df["Term"].value_counts()

Out[29]:

Short Term 68471
 Long Term 27101
 Name: Term, dtype: int64

In [30]:

1 df["Years in current job"].value_counts()

Out[30]:

10+ years 31090
 2 years 9104
 3 years 8151
 < 1 year 8114
 5 years 6778
 1 year 6436
 4 years 6132
 6 years 5676
 7 years 5573
 8 years 4569
 9 years 3949
 Name: Years in current job, dtype: int64

In [31]:

```
1 df["Home Ownership"].value_counts()
```

Out[31]:

```
Home Mortgage    46414
Rent              40550
Own Home          8403
HaveMortgage      205
Name: Home Ownership, dtype: int64
```

In [32]:

```
1 df["Purpose"].value_counts()
2
```

Out[32]:

```
Debt Consolidation    75277
other                  5665
Home Improvements     5511
Other                  3104
Business Loan         1519
Buy a Car              1208
Medical Bills         1041
Buy House              652
Take a Trip           546
major_purchase        333
small_business         272
moving                 139
wedding                111
vacation               95
Educational Expenses   89
renewable_energy       10
Name: Purpose, dtype: int64
```

In [33]:

```
1 # Label encoding the Data
2 from sklearn.preprocessing import LabelEncoder
```

In [34]:

```
1 for col in df_cat:
2     le = LabelEncoder()
3     df_cat[col]=le.fit_transform(df_cat[col])
```

In [35]:

```
1 df_cat["Loan Status"].value_counts()
```

Out[35]:

```
1    74257
0    21315
Name: Loan Status, dtype: int64
```

In [36]:

```
1 df_cat["Term"].value_counts()
```

Out[36]:

```
1    68471
0    27101
Name: Term, dtype: int64
```

In [37]:

```
1 df_cat["Years in current job"].value_counts()
```

Out[37]:

```
1    31090
2     9104
3     8151
10    8114
5     6778
0     6436
4     6132
6     5676
7     5573
8     4569
9     3949
Name: Years in current job, dtype: int64
```

In [38]:

```
1 df_cat["Home Ownership"].value_counts()
```

Out[38]:

```
1    46414
3    40550
2     8403
0      205
Name: Home Ownership, dtype: int64
```

In [39]:

```
1 df_cat["Purpose"].value_counts()
```

Out[39]:

```
3    75277
11   5665
5    5511
7    3104
0    1519
2    1208
6    1041
1     652
8     546
9     333
13    272
10    139
15    111
14     95
4      89
12     10
```


Name: Purpose, dtype: int64

In [40]:

```
1 #to check the skewness in the data
2 from scipy.stats import skew
```

In [41]:

```
1 for col in df_num:
2     print(col)
3     print(skew( df_num[col] ))
4
5     plt.figure()
6     sns.distplot(df_num[col])
7     plt.show()
```



Bankruptcies

3.5779160537124457



In [42]:

```
1 df =pd.concat([df_cat,df_num], axis = 1)
```

In [43]:

```
1 df.head()
```

Out[43]:

	Loan Status	Term	Years in current job	Home Ownership	Purpose	Current Loan Amount	Monthly Debt	Years of Credit History	Number of Open Accounts	Num of Cr Probl
0	1	1	8	1	5	445412.0	5214.74	17.2	6.0	
1	1	1	1	1	3	262328.0	33295.98	21.1	35.0	
2	1	1	8	2	3	99999999.0	29200.53	14.9	18.0	
3	1	0	3	2	3	347666.0	8741.90	12.0	9.0	
4	1	1	5	3	3	176220.0	20639.70	6.1	15.0	

In [44]:

```
1 x = df.iloc[:,1:]
2 y = df.iloc[:,0]
```

In [45]:

```
1 x
```

Out[45]:

	Term	Years in current job	Home Ownership	Purpose	Current Loan Amount	Monthly Debt	Years of Credit History	Number of Open Accounts	Number of Credit Problems
0	1	8	1	5	445412.0	5214.74	17.2	6.0	1.0
1	1	1	1	3	262328.0	33295.98	21.1	35.0	0.0
2	1	8	2	3	99999999.0	29200.53	14.9	18.0	1.0
3	0	3	2	3	347666.0	8741.90	12.0	9.0	0.0
4	1	5	3	3	176220.0	20639.70	6.1	15.0	0.0
...
99994	1	0	1	7	210584.0	3727.61	17.4	6.0	0.0
99995	1	7	2	11	147070.0	2202.86	22.3	5.0	0.0
99996	1	0	3	3	99999999.0	13109.05	9.4	22.0	0.0
99997	1	6	3	3	103136.0	7315.57	18.8	12.0	1.0
99998	1	9	3	3	530332.0	9890.07	15.0	8.0	0.0

95572 rows × 13 columns

In [46]:

```
1 y
```

Out[46]:

Loan Status	
0	1
1	1
2	1
3	1
4	1
...	...
99994	1
99995	1
99996	1
99997	1
99998	1

95572 rows × 1 columns

Model creation

In [47]:

```
1 from sklearn.model_selection import train_test_split, cross_val_score
2 xtrain, xtest, ytrain, ytest = train_test_split(x,y, test_size = 0.3, random_state =0 ,
```

In [48]:

```
1 from sklearn.ensemble import RandomForestClassifier
2 rfc = RandomForestClassifier()
3 rfc.fit(xtrain,ytrain)
4 ypred = rfc.predict(xtest)
```

In [49]:

```
1 from sklearn.metrics import classification_report, accuracy_score, confusion_matrix
```

In [50]:

```

1 print(classification_report(ytest, ypred))
2 print()
3 print(accuracy_score(ytest, ypred))
4 print()
5 print(confusion_matrix(ytest, ypred))

```

	precision	recall	f1-score	support
0	0.57	0.07	0.12	6395
1	0.79	0.99	0.87	22277
accuracy			0.78	28672
macro avg	0.68	0.53	0.50	28672
weighted avg	0.74	0.78	0.71	28672

0.7807268415178571

```

[[ 440 5955]
 [ 332 21945]]

```

In [51]:

```
1 ytrain.value_counts()
```

Out[51]:

```

Loan Status
1      51980
0      14920
dtype: int64

```

In [52]:

```
1 !pip install imblearn
```

```

Requirement already satisfied: imblearn in c:\users\arsalan\anaconda3\lib\site-packages (0.0)
Requirement already satisfied: imbalanced-learn in c:\users\arsalan\anaconda3\lib\site-packages (from imblearn) (0.8.0)
Requirement already satisfied: scipy>=0.19.1 in c:\users\arsalan\anaconda3\lib\site-packages (from imbalanced-learn->imblearn) (1.6.2)
Requirement already satisfied: scikit-learn>=0.24 in c:\users\arsalan\anaconda3\lib\site-packages (from imbalanced-learn->imblearn) (0.24.1)
Requirement already satisfied: joblib>=0.11 in c:\users\arsalan\anaconda3\lib\site-packages (from imbalanced-learn->imblearn) (1.0.1)
Requirement already satisfied: numpy>=1.13.3 in c:\users\arsalan\anaconda3\lib\site-packages (from imbalanced-learn->imblearn) (1.20.1)
Requirement already satisfied: threadpoolctl>=2.0.0 in c:\users\arsalan\anaconda3\lib\site-packages (from scikit-learn>=0.24->imbalanced-learn->imblearn) (2.1.0)

```

As the data was highly imbalanced, we need to balance the data using imblearn

In [53]:

```
1 # Synthetic minority Oversampling technique(SMOTE)
2 from imblearn.combine import SMOTETomek
3 from imblearn.under_sampling import NearMiss
```

In [54]:

```
1 smk = SMOTETomek(random_state = 42)
2 x1, y1 = smk.fit_resample(x,y)
```

In [55]:

```
1 x1.shape, y1.shape
```

Out[55]:

```
((144380, 13), (144380, 1))
```

In [56]:

```
1 from imblearn.over_sampling import RandomOverSampler
2 ros = RandomOverSampler()
3 xtrain1, ytrain1 = ros.fit_resample(x,y)
4
```

In [57]:

```
1 xtrain1.shape,ytrain1.shape
```

Out[57]:

```
((148514, 13), (148514, 1))
```

In [58]:

```
1 xtrain1, xtest, ytrain1, ytest = train_test_split(x1,y1, test_size = 0.3, random_state
```

In [59]:

```

1 from sklearn.tree import DecisionTreeClassifier
2 dt = DecisionTreeClassifier()
3 dt.fit(xtrain1,ytrain1)
4 ypred = dt.predict(xtest)
5 print(classification_report(ytest, ypred))
6 print()
7 print(accuracy_score(ytest, ypred))
8 print()
9 print(confusion_matrix(ytest, ypred))

```

	precision	recall	f1-score	support
0	0.82	0.78	0.80	21879
1	0.79	0.82	0.80	21435
accuracy			0.80	43314
macro avg	0.80	0.80	0.80	43314
weighted avg	0.80	0.80	0.80	43314

0.8023964538024657

```

[[17159  4720]
 [ 3839 17596]]

```

In [60]:

```

1 from sklearn.ensemble import RandomForestClassifier
2 rfc = RandomForestClassifier()
3 rfc.fit(xtrain1,ytrain1)
4 ypred = rfc.predict(xtest)

```

In [61]:

```

1 print(classification_report(ytest, ypred))
2 print()
3 print(accuracy_score(ytest, ypred))
4 print()
5 print(confusion_matrix(ytest, ypred))

```

	precision	recall	f1-score	support
0	0.87	0.80	0.84	21879
1	0.81	0.88	0.85	21435
accuracy			0.84	43314
macro avg	0.84	0.84	0.84	43314
weighted avg	0.84	0.84	0.84	43314

0.8415523849101907

```

[[17557  4322]
 [ 2541 18894]]

```

HyperParameter Tuning to improve the accuracy

In [62]:

```

1 rfc1 = RandomForestClassifier(criterion = "entropy")
2 rfc1.fit(xtrain1,ytrain1)
3 ypred = rfc1.predict(xtest)
4 print(classification_report(ytest, ypred))
5 print()
6 print(accuracy_score(ytest, ypred))
7 print()
8 print(confusion_matrix(ytest, ypred))

```

	precision	recall	f1-score	support
0	0.87	0.80	0.84	21879
1	0.81	0.88	0.85	21435
accuracy			0.84	43314
macro avg	0.84	0.84	0.84	43314
weighted avg	0.84	0.84	0.84	43314

0.8412753382278247

```

[[17554  4325]
 [ 2550 18885]]

```

In [63]:

```

1 rfc2 = RandomForestClassifier(max_samples = 100)
2 rfc2.fit(xtrain1,ytrain1)
3 ypred = rfc2.predict(xtest)
4 print(classification_report(ytest, ypred))
5 print()
6 print(accuracy_score(ytest, ypred))
7 print()
8 print(confusion_matrix(ytest, ypred))

```

	precision	recall	f1-score	support
0	0.67	0.71	0.69	21879
1	0.69	0.65	0.67	21435
accuracy			0.68	43314
macro avg	0.68	0.68	0.68	43314
weighted avg	0.68	0.68	0.68	43314

0.6788798079143002

```

[[15562  6317]
 [ 7592 13843]]

```

In [64]:

```

1 rfc3 = RandomForestClassifier(min_samples_leaf = 50)
2 rfc3.fit(xtrain1,ytrain1)
3 ypred = rfc3.predict(xtest)
4 print(classification_report(ytest, ypred))
5 print()
6 print(accuracy_score(ytest, ypred))
7 print()
8 print(confusion_matrix(ytest, ypred))

```

	precision	recall	f1-score	support
0	0.77	0.71	0.74	21879
1	0.73	0.79	0.76	21435
accuracy			0.75	43314
macro avg	0.75	0.75	0.75	43314
weighted avg	0.75	0.75	0.75	43314

0.7490880546705453

```

[[15531  6348]
 [ 4520 16915]]

```

In [65]:

```

1 for i in range(1,100,10):
2     rfc4 = RandomForestClassifier(min_samples_leaf = i)
3     rfc4.fit(xtrain1,ytrain1)
4     ypred = rfc4.predict(xtest)
5     print(f" {i} -: {accuracy_score(ytest, ypred)}")

```

```

1 -: 0.8421988271690447
11 -: 0.7845962044604516
21 -: 0.7694048113773837
31 -: 0.7578611996121346
41 -: 0.7579304612827261
51 -: 0.746848593988087
61 -: 0.7447245694232811
71 -: 0.7412383986701759
81 -: 0.7391143741053701
91 -: 0.7372673962229302

```

In [66]:

```
1 for i in range(1,100,10):
2     rfc5 = RandomForestClassifier(max_samples = i)
3     rfc5.fit(xtrain1,ytrain1)
4     ypred = rfc5.predict(xtest)
5     print(f" {i}    -: {accuracy_score(ytest,ypred)}")

1    -: 0.5051253636237706
11   -: 0.6271413399824537
21   -: 0.649258900124671
31   -: 0.6483815856305121
41   -: 0.6692524357020825
51   -: 0.6713302858198273
61   -: 0.6721383386433948
71   -: 0.6758553816318049
81   -: 0.6835434270674609
91   -: 0.6780948423142633
```

In [67]:

```
1 for i in range(1,1000,50):
2     rfc6 = RandomForestClassifier(n_estimators = i)
3     rfc6.fit(xtrain1,ytrain1)
4     ypred = rfc6.predict(xtest)
5     print(f" {i}    -: {accuracy_score(ytest,ypred)}")

1    -: 0.7346585399639839
51   -: 0.8372812485570486
101  -: 0.8427067460867156
151  -: 0.8438149328161795
201  -: 0.843884194486771
251  -: 0.8459851318280464
301  -: 0.8433070138985086
351  -: 0.8449692939927045
401  -: 0.844484462298564
451  -: 0.844992381216235
501  -: 0.8461005679456989
551  -: 0.8451539917809484
601  -: 0.8457311723692109
651  -: 0.8451309045574179
701  -: 0.8450847301103569
751  -: 0.8447153345338689
801  -: 0.8454310384633144
851  -: 0.8453156023456619
901  -: 0.8461698296162904
951  -: 0.8456157362515584
```

In [68]:

```
1 from sklearn.ensemble import RandomForestClassifier
2 rfc7 = RandomForestClassifier(n_estimators = 950)
3 rfc7.fit(xtrain1,ytrain1)
4 ypred = rfc7 .predict(xtest)
```

In [69]:

```

1 print(classification_report(ytest, ypred))
2 print()
3 print(accuracy_score(ytest, ypred))
4 print()
5 print(confusion_matrix(ytest, ypred))

```

	precision	recall	f1-score	support
0	0.88	0.80	0.84	21879
1	0.82	0.89	0.85	21435
accuracy			0.85	43314
macro avg	0.85	0.85	0.85	43314
weighted avg	0.85	0.85	0.85	43314

0.8455003001339059

```

[[17574 4305]
 [ 2387 19048]]

```

Using cross validation getting the mean accuracy ¶

In [70]:

```

1 cvs = cross_val_score(rfc7, x1, y1, cv = 5, scoring = "accuracy")
2 print(f"Avg. Accuracy-: {cvs.mean()}\nStandard Deviation -: {cvs.std()}")

```

Avg. Accuracy-: 0.8503670868541351
Standard Deviation -: 0.10974338801027816

RESULTS

- being a dataset based on the defaulters the percentage of defaulters are really less
- approximately the ratio of people who have paid with who has not paid is 1:4
- balancing the Dataset for an improved accuracy and f1-score
- After splitting the data into train and test we came to know the categorical data needs to be label encoded
- Using two different models we came to know Random forest classifier can be used for this data for an improved accuracy
- did some hyperparameter tuning on the model to increase the accuracy

CONCLUSION

- The best fit model for this dataset is Random forest classifier with the value of n_estimator parameter at 950
- The average accuracy of the model is 85%

In []:

```

1

```

